Implement a responsible generative AI solution in Microsoft Foundry

Thursday, March 12, 2026 4:51 PM

1. **Microsoft Responsible AI**: https://www.microsoft.com/en-us/ai/responsible-ai
2. **Map potential harms**
   1. Identify potential harms
   2. Prioritize identified harms
   3. Test and verify the prioritized harms
   4. Document and share the verified harms
3. **Measure potential harms:**
   1. Apply pre-defined criteria to evaluate the output and categorize it according to the level of potential harm it contains. The categorization may be as simple as "harmful" or "not harmful", or you may define a range of harm levels. Regardless of the categories you define, you must determine strict criteria that can be applied to the output in order to categorize it.
4. **Mitigate potential harms**
   1. Mitigation of potential harms in a generative AI solution involves a layered approach, in which mitigation techniques can be applied at each of four layers, as shown here:
      a. Model: Selecting a model that is appropriate for the intended solution use.
      b. Safety System: Microsoft Foundry includes support for content filters that apply criteria to suppress prompts and responses based on classification of content into four severity levels (safe, low, medium, and high) for four categories of potential harm (hate, sexual, violence, and self-harm).
      c. System message and grounding: Applying prompt engineering to add grounding data to input prompts, maximizing the likelihood of a relevant, nonharmful output.
      d. User experience: Designing the application user interface to constrain inputs to specific subjects or types, or applying input and output validation can mitigate the risk of potentially harmful responses.
   2. Manage a responsible generative AI solution
      a. Before releasing a generative AI solution, identify the various compliance requirements in your organization and industry and ensure the appropriate teams are given the opportunity to review the system and its documentation. Common compliance reviews include:
         i. Legal
         ii. Privacy
         iii. Security
         iv. Accessibility
      b. Implement a way for users to provide feedback and report issues. In particular, enable users to report generated content as "inaccurate", "incomplete", "harmful", "offensive", or otherwise problematic.
      c. Features in Foundry Content Safety include:
         i. Prompt shields : Scans for the risk of user input attacks on language models
         ii. Groundedness detection : Detects if text responses are grounded in a user's source content
         iii. Protected material detection : Scans for known copyrighted content
         iv. Custom categories : Define custom categories for any new or emerging patterns
5. **Exercise**: https://microsoftlearning.github.io/mslearn-ai-studio/Instructions/Exercises/06-Explore-content-filters.html (Use this link to know how to create guardrails in Foundry)
6. Module Assessment:
   1.
   2. 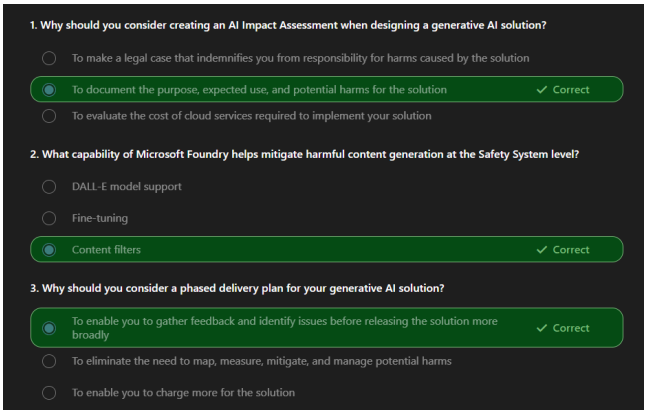

   3.